
When Actions Disappear: Adversarial Action Removal in Self-Play Reinforcement Learning

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 We study adversarial action masking in self-play reinforcement learning: an at-
2 tacker selectively removes legal actions from a victim’s action set. Unlike ob-
3 servation or action perturbations, removal eliminates decision options before the
4 agent acts. Across poker games scaling from 6 to 5,531 information states and
5 two non-poker domains, learned masking causes substantially more damage than
6 random masking and learned perturbation baselines. The attack persists across
7 Q-learning, PPO, NFSP, neural NFSP, and DQN victims; transfers across agents;
8 is amplified by self-play; and shows no recovery under extended masked train-
9 ing. Mechanistically, the adversary targets high-value decision points, captured
10 by reach-weighted contingent action capacity (CAC_w) and a value-weighted re-
11 finement CAC_v . These results identify action availability as a distinct robustness
12 surface in self-play RL.

13 1 Introduction

14 Multi-agent reinforcement learning (MARL) agents trained via self-play achieve strong performance
15 in competitive domains [Silver et al., 2018, Brown and Sandholm, 2019], but their robustness to
16 structural environment changes remains poorly understood. Prior work on adversarial attacks fo-
17 cuses on observation perturbations [Huang et al., 2017, Gleave et al., 2020] or reward manipula-
18 tion [Zhang et al., 2020]. We study a different, more severe attack surface: the *action space* itself.

19 An adversary that selectively removes actions—disabling specific capabilities rather than adding
20 noise—poses a qualitatively different threat from bounded perturbation. Such attacks arise naturally:
21 hardware failures disable actuators, regulatory changes restrict strategies, API deprecations remove
22 endpoints, and sandboxing constrains agent capabilities. We formalise this as a bi-level optimisation
23 where the inner loop trains an RL agent under masked actions and the outer loop trains an adversary
24 to choose which actions to remove.

25 Our key finding is that adversarial masking is *dramatically more efficient* than random removal:
26 across Leduc variants, learned removal is up to $4.8\times$ more damaging than random masking at com-
27 parable support. The mechanism operates through selective minimisation of reach-weighted contin-
28 gent action capacity (CAC_w)—empirically verified by strong correlation between CAC_w and victim
29 reward across budget levels.

30 Contributions.

- 31 • We formalise adversarial action masking and show it is $4\times$ more damaging than learned
32 perturbation (RARL-style) with equal training budget.

- We demonstrate scaling from 6 to **5,531** information states across five algorithms (QL, PPO, NFSP, neural NFSP, DQN), with the adversary’s advantage increasing from $2.2\times$ to $4.8\times$ within the Leduc family as game complexity grows.
- We validate cross-domain in two non-poker environments (competitive gridworld, resource collection), confirming the phenomenon is not poker-specific.
- We connect the mechanism to CAC_w ($r = 0.80$) and a refined CAC_v ($r = 0.81$), and show victims do not recover under extended training.

2 Related Work

Adversarial attacks on RL. Adversarial attacks on RL most commonly perturb observations [Huang et al., 2017, Gleave et al., 2020, Zhang et al., 2021, Sun et al., 2022] or poison rewards [Zhang et al., 2020]. In multi-agent settings, recent work studies adversarial policies [Gleave et al., 2020], sparse or model-based attacks against cooperative MARL [Lin et al., 2020, Hu and Zhang, 2022], and collusive policy-level attackers [Niu et al., 2026]. These attacks manipulate observations, policies, or rewards while leaving the legal action set intact. We instead study a structural attack on action availability itself.

Robust and action-robust RL. Robust MDPs [Iyengar, 2005, Nilim and El Ghaoui, 2005] and constrained MDPs [Altman, 1999] model uncertainty or constraints in the environment. Robust Adversarial RL (RARL) [Pinto et al., 2017] and action-robust RL [Tessler et al., 2019] train against perturbations of the action actually executed. Those channels modify selected actions; our adversary removes actions from the legal set before selection. This difference matters: bounded perturbations preserve the agent’s ability to choose, whereas removal can collapse a decision point to a singleton action set.

Robust MARL defenses. Recent robust MARL methods train agents against structured adversaries, including game-theoretic robust training for temporally coupled perturbations (GRAD) [Liang et al., 2024] and fault-switching MARL defenses such as MARTA [Mguni et al., 2025]. These approaches are complementary: they improve robustness to perturbations, malfunctions, or adversarial agents within an intact joint action space. Our attack changes the feasible action set itself. A mask-aware version of such defenses is a natural direction, but existing guarantees do not directly cover state-dependent removal of legal actions.

Action sets and masking. Invalid action masking prevents agents from selecting impossible moves [Huang and Ontañón, 2022]. Decision-theoretic planning has long studied structural leverage and action contingencies [Boutilier et al., 1999], but typically not adversarially selected action availability during learning. Our work reverses the usual masking motivation: the mask is not a safety or efficiency aid, but an attacker-chosen capability removal.

Self-play and imperfect-information games. Self-play can overfit, cycle, or exploit non-transitive structure [Balduzzi et al., 2019, Lanctot et al., 2019]. Regret-minimization methods such as CFR [Zinkevich et al., 2007] and Deep CFR [Brown et al., 2019] provide strong baselines for imperfect-information games, while NFSP [Heinrich and Silver, 2016] combines best-response learning with average-strategy tracking. Population methods such as PSRO [Lanctot et al., 2017] maintain diversity. We show that these algorithmic stabilizers do not address a different failure mode: when the action set itself is structurally reduced, averaging and population diversity cannot restore eliminated strategic dimensions.

3 Problem Formulation

Game model. Consider a two-player zero-sum extensive-form game Γ with information sets $\mathcal{I} = \mathcal{I}_0 \cup \mathcal{I}_1$. At information set $h \in \mathcal{I}_p$, player p chooses from legal actions $A(h)$. Player 0 is the *victim*; Player 1 is the opponent. Both learn via self-play.

Adversary definition. An *action-removal adversary* is a mapping $\mathcal{M} : \mathcal{I}_0 \rightarrow 2^A$ that, at each victim information set h , selects a subset $\mathcal{M}(h) \subseteq A(h)$ of actions to *retain*. The victim observes only $\mathcal{M}(h)$ and is unaware of the adversary. Formally:

- **Input:** information set h , legal actions $A(h)$, current player p .
- **Output:** $\mathcal{M}(h) \subseteq A(h)$ with $|\mathcal{M}(h)| \geq 1$ (at least one action retained).
- **Constraint:** $|\text{supp}(\mathcal{M})| = |\{h : |\mathcal{M}(h)| < |A(h)|\}| \leq k$ (budget).
- **Objective:** minimise the victim’s expected value $V_0(\pi_{\mathcal{M}}^*)$.

Bi-level optimisation. The adversary and victim interact through a bi-level problem:

$$\text{Inner: } \pi_{\mathcal{M}}^* = \arg \max_{\pi} \mathbb{E} \left[\sum_t r_t \mid \pi, \mathcal{M} \right] \quad (1)$$

$$\text{Outer: } \mathcal{M}^* = \arg \min_{\mathcal{M} \in \mathcal{C}_k} V_0(\pi_{\mathcal{M}}^*) \quad (2)$$

where $\mathcal{C}_k = \{\mathcal{M} : |\text{supp}(\mathcal{M})| \leq k\}$. The inner loop trains the victim under the adversary’s mask via RL; the outer loop updates the adversary via REINFORCE with reward signal $-V_0$. In practice, we alternate: 500 episodes of inner training, then one adversary gradient step, for 20–25 outer iterations.

Adversary implementation. *Tabular:* preference table $\theta(h, a)$ with $p_{\text{remove}}(a|h) = \text{softmax}(\theta(h, \cdot))$; the action with highest removal probability is removed. Updated via REINFORCE. *Neural:* MLP $f_{\phi} : \mathbb{R}^d \rightarrow \Delta^{|A|+1}$ mapping state features to a distribution over $|A|$ removal choices plus “no removal.” Trained via REINFORCE with mean-reward baseline.

Connection to CAC_w . Define *reach-weighted contingent action capacity*:

$$\text{CAC}_w(\mathcal{M}) = \sum_{h \in \mathcal{I}_0} \rho(h) \cdot \mathbf{1}[|\mathcal{M}(h)| > 1]$$

where $\rho(h)$ is the reach probability of h under current play. Each masked state reduces CAC_w . The adversary’s value-minimisation objective (2) can be decomposed: masking state h reduces V_0 by approximately $\rho(h) \cdot \delta(h)$, where $\delta(h) = |Q(h, a^*) - Q(h, a_{\text{forced}})|$ is the *value gap* at h . This motivates a refined measure:

$$\text{CAC}_v(\mathcal{M}) = \sum_{h \in \mathcal{I}_0} \rho(h) \cdot \delta(h) \cdot \mathbf{1}[|\mathcal{M}(h)| > 1]$$

Empirically, CAC_v correlates with victim reward at $r = 0.81$ vs. CAC_w ’s $r = 0.77$ (Section 5.4).

Deterministic Exploitation Attractor (DEA).

Proposition 1 (DEA Convergence). *Under $\text{CAC}_w = 0$ in self-play Q-learning with ε -greedy ($\varepsilon > 0$, $\alpha \in (0, 1)$): (i) the victim’s policy converges to the unique forced action at every information set; (ii) the opponent’s Q-values converge to $Q^* = V^{BR(\sigma^*)}$; (iii) $(\sigma^*, BR(\sigma^*))$ is a stable fixed point.*

Proof sketch. Forced actions make the victim’s policy Q-value-independent. The opponent faces a stationary MDP; Q-learning converges under standard conditions. The pair is stable: neither player can unilaterally deviate.

Proposition 2 (Damage bound). *Let $\mathcal{M}$ mask a single information set h^* by retaining a singleton action $a_f(h^*)$. The victim’s value loss relative to unmasked play is bounded by:*

$$\Delta V_0 \geq \rho(h^*) \cdot [Q_0(h^*, a^*(h^*)) - Q_0(h^*, a_f(h^*))]$$

where $a^*(h) = \arg \max_{a \in A(h)} Q_0(h, a)$ is the victim’s best action, $a_f(h)$ is the action forced by the mask at h , and $\rho(h)$ is the reach probability. For a budget- k adversary with support $S = \text{supp}(\mathcal{M})$, the additive approximation is:

$$\Delta V_0(\mathcal{M}) \approx \sum_{h \in S} \rho(h) \cdot [Q_0(h, a^*(h)) - Q_0(h, a_f(h))]$$

112 and is upper-bounded by:

$$\Delta V_0(\mathcal{M}) \leq \sum_{h \in S} \rho(h) \cdot \max_{a_f \in A(h) \setminus \{a^*(h)\}} [Q_0(h, a^*(h)) - Q_0(h, a_f)].$$

113 Thus a greedy adversary selects states with high $\rho(h)\delta(h)$, where $\delta(h) = Q_0(h, a^*(h)) -$
 114 $Q_0(h, a_f(h))$ —exactly the states that maximise per-unit CAC_v reduction.

115 *Proof.* The bound follows from the definition of Q-values and linearity of expectation over the
 116 reach-weighted trajectory. At state h^* , forcing $a_f(h^*)$ instead of $a^*(h^*)$ loses $Q_0(h^*, a^*(h^*)) -$
 117 $Q_0(h^*, a_f(h^*))$ in expectation, weighted by $\rho(h^*)$. Under the simplifying assumption that masked
 118 states lie on independent branches, the resulting coverage objective suggests an approximate greedy
 119 structure.¹

120 **Theoretical scope.** Propositions 1–2 provide *sufficient conditions* and *local bounds*, not a com-
 121 plete characterization of the optimal adversary. The damage bound is per-state and additive, ignoring
 122 cross-state interactions (masking state h may shift reach probabilities and Q-values at downstream
 123 states). A global characterization would require solving a combinatorial optimisation over $\binom{|Z_0|}{k}$
 124 mask configurations—NP-hard in general via reduction to weighted max coverage. Under approxi-
 125 mate independence, the value-gap decomposition motivates a greedy heuristic over high $\rho \cdot \delta$ states;
 126 we do not claim a general approximation guarantee for arbitrary extensive-form games. We view
 127 these bounds as explaining *why* the adversary works rather than fully characterizing *how well* it can
 128 work.

129 4 Methods

130 **Self-play protocol.** Both players share a single learning agent; the mask is applied to Player 0’s
 131 actions at every decision point during *both training and evaluation*. In the self-play regime, both
 132 sides adapt continuously. In the fixed-opponent regime, a snapshot of the agent’s value function
 133 is frozen before masking begins; this static copy plays as Player 1 while only Player 0 continues
 134 learning under the mask.

135 **Victim agents.** We evaluate four algorithms spanning tabular and neural methods:

- 136 • **Tabular Q-Learning:** ε -greedy, $\varepsilon = 0.15$, $\alpha = 0.1$.
- 137 • **Tabular PPO:** softmax policy, clipped surrogate, entropy bonus (Appendix A).
- 138 • **Tabular NFSP** [Heinrich and Silver, 2016]: best-response QL + average strategy, $\eta = 0.1$.
- 139 • **DQN:** 2-layer MLP (64 hidden), experience replay (20k buffer), target network, Adam
 140 (10^{-3}) . Full hyperparameters in Appendix C.

Table 1: Masking strategies evaluated.

Strategy	Description
None	Full action space (baseline)
Random(p)	Remove each action independently with prob. p
Fixed	Always remove a specific action (e.g., RAISE)
Adversarial (tabular)	Learned per-state removal via softmax preference table
Adversarial (neural)	MLP adversary trained with REINFORCE
Value heuristic	Remove at states with highest $ Q $

141 Masking strategies.

142 **Adversary architectures.** Both tabular and neural adversaries are defined formally in Section 3.
 143 The neural adversary scales its architecture with game size: 2-layer (32 hidden) for Kuhn/Leduc,
 144 3-layer (128-64) for Leduc-10/20.

¹This independence assumption is only approximate in extensive-form games: masking one state can alter reach probabilities and values downstream.

145 **Environments.** **Kuhn Poker:** 3 cards, 2 actions, 6 P0 info states. **Leduc Poker:** 6 cards (3 ranks
146 $\times$ 2 suits), 3 actions, 2 rounds, ~ 50 P0 info states. **Leduc-5:** 10 cards (5 ranks), 389 P0 info states.
147 **Leduc-10:** 20 cards (10 ranks), 1,496 P0 info states. **Leduc-20:** 40 cards (20 ranks), **5,531 P0**
148 **info states.** **Competitive Gridworld:** 5×5 , 5 actions, 149 P0 states. **Resource Collection:** 4×4 , 4
149 actions, 18,604 P0 states. All implemented from scratch with MaskedEnv wrappers.

150 5 Experiments

151 We structure experiments around five questions: (1) Does adversarial masking scale with game
152 complexity? (2) Is the vulnerability algorithm-invariant? (3) Does it generalise beyond poker?
153 (4) How efficient is adversarial targeting? (5) What mechanism drives the attack?

154 5.1 Scaling with Game Complexity

155 Our strongest result uses function approximation on both sides: a DQN victim and a neural MLP
156 adversary in Leduc variants whose state spaces range from ~ 50 to 5,531 victim information states.




Figure 1: Scaling trend. The adversarial/random damage ratio is plotted against victim information-set count on a log scale. Error bars show 95% CIs from the DQN scale experiments.

157 We implement Leduc-N, parameterised by rank count. **Leduc-5** (10 cards, 389 states), **Leduc-10** (20
158 cards, 1,496 states), and **Leduc-20** (40 cards, **5,531 states**) test whether the adversary’s advantage
159 persists as game complexity grows by $920\times$ from Kuhn. Figure 1 and Table 2 show a consistent
160 empirical scaling trend.

Table 2: DQN victim across poker scales (5 seeds, 95% CIs).

Game	$ I_0 $	None	Random	Neural Adv	Ratio
Leduc	~ 50	-0.10	-1.17	-2.57 ± 0.47	$2.2\times$
Leduc-5	389	-0.15	-0.60	-2.74 ± 0.23	$4.6\times$
Leduc-10	1,496	-0.18	-0.68	-3.19 ± 0.23	$4.7\times$
Leduc-20	5,531	-0.09	-0.63	-3.00 ± 0.15	$4.8\times$

161 At 5,531 states the neural adversary drives DQN to -3.00 per hand— $4.8\times$ worse than random—
162 with the tightest CIs of any scale (± 0.15). A strict matched- L_0 control in Leduc samples exactly
163 the same number of states as the learned mask ($k = 64.8 \pm 4.3$): adversarial masking reaches
164 -2.32 ± 0.36 while matched random reaches -1.03 ± 0.24 ($2.24\times$ gap). The advantage is therefore
165 *which* states are targeted, not support size.

166 **Threat-model ablation.** The main setting assumes environment-level control: the adversary can
167 observe the victim information state and alter legal actions before selection. This is intentionally

strong, matching settings where a simulator, platform, or operating layer controls capabilities. To test whether private-card access is essential, we restrict the adversary to public information only (betting history and public card, private rank hidden). In Leduc, public-info adversarial masking still reaches -1.71 ± 0.58 , substantially worse than random masking (-0.98 ± 0.08), though weaker than private-info masking (-2.30 ± 0.39). Thus strong private observability is not required: public structural information is sufficient for damaging targeted removal.

Neural NFSP on Leduc-5. To confirm the result holds with a stronger victim, we run *Neural NFSP* (128-64 MLP average policy, DQN best-response, reservoir-sampled action buffer) on Leduc-5. Pre-attack: -0.13 ± 0.09 ; post-adversarial: -1.89 ± 0.16 ; post-random: -0.89 ± 0.06 . The adversary is $2.3\times$ more damaging than random, confirming that even proper neural game-solving methods collapse under adversarial action removal.




Figure 2: Learning under a fixed mask. NFSP in Leduc-5 does not recover under continued masked training; adversarial removal remains substantially worse than random removal.

DQN masking mechanics. Masked actions are filtered from the legal action list *before* ϵ -greedy selection: the DQN computes Q-values for all actions, then restricts argmax and exploration to the masked subset. The victim is unaware of the masking. Total DQN training: 20k pre-training episodes (no mask) + 20 outer $\times$ 500 inner = 10k episodes under the mask. Learning curves (Appendix F) confirm convergence: reward stabilises by episode $\sim 8k$ under the mask.

OpenSpiel/CFR context. Kuhn and Leduc rules match the standard imperfect-information games used in OpenSpiel [Lanctot et al., 2019]; we implement them from scratch only to expose the action-mask hook at every legal-action query and to scale Leduc-N by rank count. CFR and Deep CFR [Zinkevich et al., 2007, Brown et al., 2019] assume a fixed game tree. Action removal changes the game tree itself, so CFR remains a useful equilibrium reference but is not a direct defense unless retrained on the masked game. Our results should therefore be read as structural robustness tests of learning dynamics under game-tree modification.

5.2 Tabular Validation Across Algorithms

All algorithms collapse under adversarial masking—in both Kuhn and Leduc. The vulnerability is *algorithm-invariant*: Q-Learning, PPO, and NFSP all degrade to comparable levels. NFSP on Leduc drops from -0.01 to -1.58 (raw), confirming that NFSP’s average-strategy component cannot compensate when the action space itself is reduced. Combined with the DQN scaling result (Table 2), we have **four algorithms across two games and two paradigms (tabular + neural) all exhibiting collapse**.

Table 3: Tabular agents in Kuhn and Leduc under masking (5 seeds). Both raw reward (r) and normalised performance shown.

Setting	Raw Reward		Normalised	
	None	Advers.	None	Advers.
Kuhn + QL	−0.05	−0.98	.487	.267
Kuhn + PPO	−0.05	−0.85	.486	.274
Kuhn + NFSP	+0.12	−0.45	.530	.388
Leduc + QL	+0.05	−3.50	.502	.367
Leduc + PPO	−0.08	−3.32	.497	.369
Leduc + NFSP	−0.01	−1.58	.500	.440

5.3 Cross-Domain Validation: Competitive Gridworld

To verify the phenomenon is not a poker artifact, we test on a *completely different* domain: a 5×5 competitive gridworld where Player 0 (prey) navigates to a goal while Player 1 (predator) attempts to intercept. Both players have 5 actions (UP/DOWN/LEFT/RIGHT/STAY), the game is turn-based with perfect information, and 149 unique P0 states are observed.

Table 4: Cross-domain: Competitive Gridworld (5×5 , 5 actions, 149 P0 states, 5 seeds).

Masking	P0 Reward	Δ
None	$+0.70 \pm 0.01$	—
Random ($p=0.3$)	$+0.04 \pm 0.18$	−0.66
Fixed (UP)	$+0.71 \pm 0.10$	+0.01
Adversarial	$−0.58 \pm 0.03$	−1.27

The adversary causes $1.9 \times$ more damage than random masking in a domain with no cards, no imperfect information, and a spatial rather than strategic structure. Fixed removal (always remove UP) has negligible effect because the prey reroutes; the adversary *state-dependently* removes the action that matters most at each position.

Resource Collection (4×4). A second non-poker domain: two agents compete to collect 4 resources on a grid, with reward = resources_collected_difference. 4 actions, 18,604 unique P0 states observed. Adversarial masking degrades P0 by $\Delta = -0.13$ vs. random $\Delta = -0.10$ ($1.4 \times$). Fixed removal (UP) *helps* P0 ($\Delta = +0.32$), showing that only state-dependent targeting produces consistent harm. The smaller gap ($1.4 \times$ vs. $1.9 \times$) reflects the symmetric, coordination-like structure where many states have similar strategic value—there is less exploitable structure to target.

Two non-poker domains with different structures (goal-seeking vs. resource competition, perfect information, spatial) both show adversarial $>$ random. The phenomenon is not poker-specific.

5.4 Attack Efficiency and CAC_w Correlation

Budget sweep. Appendix E shows victim reward as a function of adversary budget k (number of info states masked). Targeted removal degrades the victim as support grows; random masking requires broad coverage to approach the same damage.

CAC_w correlation. To verify that the adversary operates through CAC_w minimisation, we compute reach-weighted CAC_w at each budget level and correlate with victim reward. Across all budget levels and both adversarial and random conditions (10 seeds each), the Pearson correlation between CAC_w and reward is $r = 0.77$ ($p < 0.001$); for adversarial-only conditions, $r = 0.80$. The full budget table is in Appendix E.

CAC_v -greedy oracle. We also test a direct oracle heuristic that selects the top- k Kuhn states by estimated $\rho(h)\delta(h)$ and removes the victim’s estimated best action. At $k = 3$, CAC_v -greedy reaches -0.29 ± 0.09 , stronger than random (-0.19 ± 0.06) and the learned adversary under the same short

227 training budget (-0.22 ± 0.02). This supports CACv as a mechanistic proxy: when given the value-
 228 gap statistic directly, greedy targeting produces the expected harm.

229 **Estimating CACv.** We estimate $\rho(h)$ from on-policy visitation frequencies during evaluation roll-
 230 outs and $\delta(h)$ from the victim’s learned Q-table or network Q-values at the same checkpoint. Thus
 231 CACv is a diagnostic proxy, not an oracle ground truth: Q-estimation error can perturb rankings.
 232 Despite this estimation noise, CACv rankings are stable across seeds and align with oracle perfor-
 233 mance: the top CACv states in Kuhn repeatedly include the same high-impact information sets (0pb,
 234 1, 2pb), and the CACv-greedy ablation remains stronger than random.

235 5.5 Action Removal vs. Perturbation (RARL)

Table 5: Action removal vs. perturbation: both with *learned* adversary (Kuhn, 10 seeds, 95% CIs). Both adversaries use the same REINFORCE training (20 outer $\times$ 500 inner). Removal is $4\times$ more damaging even when the perturbation adversary is given equal training budget.

Attack Type	Raw Reward	Normalised
None	-0.02 ± 0.07	.494
Learned perturbation	-0.24 ± 0.09	.440
Learned removal	-1.01 ± 0.09	.248

236 Both adversaries use identical training (REINFORCE, same budget, same iterations), ensuring a fair
 237 comparison. The learned perturbation adversary (-0.24) is stronger than fixed- p RARL (-0.09 , not
 238 shown), but learned removal (-1.01) remains $4\times$ more damaging. The mechanism is qualitatively
 239 different: perturbation adds noise to execution but preserves the agent’s *ability to choose*; removal
 240 eliminates decision points entirely, reducing CAC_w [Pinto et al., 2017, Tessler et al., 2019].

241 **Perturbation protocol.** In the learned perturbation baseline, the adversary observes the same in-
 242 formation state and replaces the victim’s chosen action with another legal action (or no-op) before
 243 execution. The victim is unaware of the replacement and receives the same training/evaluation
 244 schedule as in the removal experiment. This matches training budget and observability while pre-
 245 serving the legal action set.

246 5.6 Additional Ablations

247 The attack is amplified by self-play co-adaptation (self-play: -0.975 , fixed opponent: -0.841),
 248 transfers across random seeds (transferred: -1.032 vs. per-agent retrained: -0.872), transfers from
 249 Q-learning to NFSP ($\Delta = -0.49$), and is not an artifact of parameter sharing: separate tabu-
 250 lar Q-tables match shared-agent results, while separate DQN networks in Leduc still collapse to
 251 -1.53 ± 0.48 . Evaluation-only masking already degrades a normally trained Leduc victim (-0.58),
 252 while continued and mask-aware training converge to larger losses (-2.65 and -2.71), showing
 253 both immediate structural harm and long-run adaptation failure. Simple action-dropout helps mod-
 254 estly (-1.82 vs. -2.45), but a persistent random mask-ensemble defense does not help (-2.64).
 255 This negative result is informative: robustness to arbitrary stochastic unavailability is not enough;
 256 effective defenses likely need targeted protection or redundancy at high-CACv states. Full ablation
 257 tables are in Appendix E.

258 6 Discussion

259 **Scaling trend.** The scaling trajectory (Table 2) shows a consistent empirical trend. Excluding
 260 Kuhn (2-action ceiling effect), log-linear regression of adversary damage ratio on $\log_{10}(|\mathcal{I}_0|)$ yields
 261 $R^2 = 0.62$ (Pearson $r = 0.79$, $n = 4$ game sizes), with slope 1.56. We report this as an *empirical*
 262 *trend*, not a scaling law—four data points establish direction, not a precise functional form. The
 263 absolute damage gap (adversarial minus random reward) grows from -1.40 at 50 states to -2.51
 264 at 1,496 states, and the adversary’s variance *decreases* with scale (CI $\pm 0.47 \rightarrow \pm 0.15$), suggesting
 265 more robust convergence in larger games where there is more targetable strategic heterogeneity.
 266 Neural NFSP ($2.3\times$ on Leduc-5) confirms this holds for properly neural game-solving methods.

CACw: necessary but not sufficient. The CAC_w -reward correlation ($r = 0.80$) is consistent with the hypothesis that reducing strategic flexibility drives performance degradation. However, CAC_w is a *necessary* condition for collapse, not a complete characterization: at intermediate budgets, the adversary sometimes achieves comparable or greater damage with higher CAC_w than random masking (Table 19, $k=3$: adversarial $CAC_w = 0.68$ vs. random 0.37, yet adversarial reward = -0.25 vs. random -0.21). This occurs because CAC_w aggregates over all reachable states; the adversary targets the *strategically pivotal* subset where action removal most disrupts equilibrium play, whereas random removal may eliminate more states but at less important ones. We test a refined measure, *counterfactual-value-weighted capacity* $CAC_v = \sum_h \rho(h) \cdot |Q(h, a_0) - Q(h, a_1)| \cdot \mathbf{1}[|\mathcal{M}(h)| > 1]$, which weights each state by its Q-value gap (strategic importance). CAC_v achieves Pearson $r = 0.81$ vs. CAC_w 's $r = 0.77$, confirming that incorporating value sensitivity improves the structural explanation. The remaining unexplained variance likely reflects higher-order interactions (e.g., how masking one state changes the value landscape at downstream states).

Implications for deployment. Real-world multi-agent systems where an adversary can disable specific capabilities (API endpoints, actuators, communication channels) are vulnerable to targeted collapse. Defences should maintain strategic flexibility at high-reach, high- CAC_v decision points. Uniform robustness methods such as random action dropout or random mask ensembles are insufficient in our experiments; the defensive object is not raw action count, but redundancy at the strategically pivotal states where removal produces large value gaps. Designing effective defenses likely requires identifying and preserving high- CAC_v states, rather than optimizing for uniform robustness to arbitrary action unavailability.

7 Limitations

Our largest game (Leduc-20) has $\sim 5,500$ information states—nearly three orders of magnitude larger than Kuhn. The increasing adversary advantage with scale (Table 2) suggests the mechanism strengthens as games grow, but verification at full poker scale (10^5+ states) remains future work.

Discrete scope. Action removal is inherently a *discrete* attack: it eliminates specific actions from a finite set. Our results—across poker, gridworld, and resource collection—all use small discrete action spaces ($|A| \leq 5$). In continuous-action domains (e.g., MuJoCo), the analogous attack would be *region exclusion* (disabling subsets of the action manifold), which requires different formalisation and may exhibit qualitatively different dynamics. We do not claim generality beyond discrete action spaces.

Theory. The damage bound (Proposition 2) is local and additive, assuming independent masking effects across states. Cross-state interactions (masking one state shifts Q-values at others) are not captured. The scaling regression ($R^2 = 0.62$, 4 game sizes) establishes a trend but not a precise functional relationship.

8 Conclusion

We show that self-play RL agents—from tabular Q-Learning to neural NFSP and DQN, from 6-state Kuhn to 5,531-state Leduc-20, across poker and non-poker domains—are brittle to targeted action-space attacks. Within the Leduc family, the adversary's advantage over random masking increases with game complexity ($2.2\times$ at 50 states to $4.8\times$ at 5,531), and the vulnerability persists across five algorithms and seven environments. The mechanism operates through selective reduction of strategically important decision capacity (CAC_w/v), victims do not recover under extended training, and the attack transfers across agents. These findings argue for robustness designs that preserve strategic flexibility at high-reach decision points rather than raw action count.

References

Eitan Altman. *Constrained Markov decision processes: Stochastic modeling*. CRC Press, 1999.

David Balduzzi, Marta Garnelo, Yoram Bachrach, Wojciech Czarnecki, Julien Pérolat, Max Jaderberg, and Thore Graepel. Open-ended learning in symmetric zero-sum games. In *International Conference on Machine Learning*, pages 434–443, 2019.

Craig Boutilier, Thomas Dean, and Steve Hanks. Decision-theoretic planning: Structural assumptions and computational leverage. *Journal of Artificial Intelligence Research*, 11:1–94, 1999.

Noam Brown and Tuomas Sandholm. Superhuman AI for multiplayer poker. *Science*, 365(6456):885–890, 2019.

Noam Brown, Adam Lerer, Sam Gross, and Tuomas Sandholm. Deep counterfactual regret minimization. In *International Conference on Machine Learning*, pages 793–802, 2019.

Adam Gleave, Michael Dennis, Cody Wild, Neel Kant, Sergey Levine, and Stuart Russell. Adversarial policies: Attacking deep reinforcement learning. In *International Conference on Learning Representations*, 2020.

Johannes Heinrich and David Silver. Deep reinforcement learning from self-play in imperfect-information games. In *International Conference on Learning Representations*, 2016.

Yizheng Hu and Zhihua Zhang. Sparse adversarial attack in multi-agent reinforcement learning. *arXiv preprint arXiv:2205.09362*, 2022.

Sandy Huang, Nicolas Papernot, Ian Goodfellow, Yan Duan, and Pieter Abbeel. Adversarial attacks on neural network policies. *arXiv preprint arXiv:1702.02284*, 2017.

Shengyi Huang and Santiago Ontañón. A closer look at invalid action masking in policy gradient algorithms. In *International FLAIRS Conference*, 2022.

Garud N. Iyengar. Robust dynamic programming. *Mathematics of Operations Research*, 30(2):257–280, 2005.

Marc Lanctot, Vinicius Zambaldi, Audrunas Gruslys, Angeliki Lazaridou, Karl Tuyls, Julien Pérolat, David Silver, and Thore Graepel. A unified game-theoretic approach to multiagent reinforcement learning. In *Advances in Neural Information Processing Systems*, volume 30, 2017.

Marc Lanctot, Edward Lockhart, Jean-Baptiste Lespiau, Vinicius Zambaldi, et al. OpenSpiel: A framework for reinforcement learning in games. *arXiv preprint arXiv:1908.09453*, 2019.

Yongyuan Liang, Yanchao Sun, Ruijie Zheng, Xiangyu Liu, Benjamin Eysenbach, Tuomas Sandholm, Furong Huang, and Stephen McAleer. Game-theoretic robust reinforcement learning handles temporally-coupled perturbations. In *International Conference on Learning Representations*, 2024.

Jieyu Lin, Kristina Dzevaroska, Sai Qian Zhang, Alberto Leon-Garcia, and Nicolas Papernot. On the robustness of cooperative multi-agent reinforcement learning. In *IEEE Security and Privacy Workshops*, pages 62–68, 2020.

David Mguni, Yaqi Sun, Haojun Chen, Wanrong Yang, Amir Darabi, Larry Olanrewaju Orimoloye, and Yaodong Yang. Learning robust multi-agent policies via selective adversarial fault induction. *arXiv preprint arXiv:2508.08800*, 2025.

Arnab Nilim and Laurent El Ghaoui. Robust control of Markov decision processes with uncertain transition matrices. *Operations Research*, 53(5):780–798, 2005.

Men Niu, Xinxin Fan, Quanliang Jing, Shaoye Luo, and Yunfeng Lu. CAMA: Exploring collusive adversarial attacks in cooperative multi-agent reinforcement learning. *arXiv preprint arXiv:2603.20390*, 2026.

Lerrel Pinto, James Davidson, Rahul Sukthankar, and Abhinav Gupta. Robust adversarial reinforcement learning. In *International Conference on Machine Learning*, pages 2817–2826, 2017.

David Silver, Thomas Hubert, Julian Schrittwieser, et al. A general reinforcement learning algorithm that masters chess, shogi, and Go through self-play. *Science*, 362(6419):1140–1144, 2018.

- 360 Yanchao Sun, Ruijie Zheng, Yongyuan Liang, and Furong Huang. Who is the strongest enemy?
361 towards optimal and efficient evasion attacks in deep RL. *International Conference on Learning*
362 *Representations*, 2022.
- 363 Chen Tessler, Yonathan Efroni, and Shie Mannor. Action robust reinforcement learning and appli-
364 cations in continuous control. *International Conference on Machine Learning*, pages 6215–6224,
365 2019.
- 366 Huan Zhang, Hongge Chen, Duane Boning, and Cho-Jui Hsieh. Robust reinforcement learning on
367 state observations with learned optimal adversary. *International Conference on Learning Repre-*
368 *sentations*, 2021.
- 369 Xuezhou Zhang, Yuzhe Ma, Adish Singla, and Xiaojin Zhu. Adaptive reward-poisoning attacks
370 against reinforcement learning. *International Conference on Machine Learning*, 2020.
- 371 Martin Zinkevich, Michael Johanson, Michael Bowling, and Carmelo Piccione. Regret minimization
372 in games with incomplete information. In *Advances in Neural Information Processing Systems*,
373 volume 20, 2007.

374 Reproducibility Statement

375 All environments, agents, adversaries, and experiment scripts are implemented in the accompanying
376 repository. Each reported result is generated by a standalone script under `experiments/`; random
377 seeds are fixed in the scripts, and hyperparameters are listed in Appendix C. Figures are generated by
378 `experiments/generate_neurips_figures.py` from completed experiment outputs. The largest
379 run, Leduc-20, uses five seeds, 30k victim pre-training episodes, and 25 adversary outer iterations
380 with 500 inner episodes each.

381 A Tabular PPO Details

382 Our tabular PPO maintains a softmax policy over a preference table $\theta(s, a)$. Action prob-
383 abilities: $\pi(a|s) = \text{softmax}(\theta(s, \cdot))$. Updates use the clipped surrogate objective $L =$
384 $\min(r_t A_t, \text{clip}(r_t, 1-\epsilon, 1+\epsilon) A_t)$ with $r_t = \pi_{\text{new}}/\pi_{\text{old}}$, advantage $A_t = R - b$ (running baseline),
385 clip parameter $\epsilon = 0.2$, entropy bonus coefficient 0.01, learning rate 0.01.

386 B Experimental Protocols

387 **Self-play protocol.** Both victim and opponent are the same agent; the mask is applied to Player 0’s
388 actions at every decision during both training and evaluation. In the fixed-opponent regime, a snap-
389 shot of the Q-table is frozen before masking; this static copy serves as Player 1 while only Player 0
390 continues learning.

391 **Transfer experiment.** The adversary is trained against one victim (seed 42) for 20 outer $\times$ 500
392 inner episodes. This mask is applied without modification to fresh victims (seeds 123–2048), each
393 pre-trained 10k episodes then trained 10k under the transferred mask.

394 **Budget enforcement.** Budget k is enforced via top- k projection: the adversary trains over all
395 states but at execution only the top- k by confidence (max softmax prob. minus uniform baseline)
396 have masks applied. This approximates a hard L_0 constraint on mask support.

397 **Action-removal granularity.** Unless otherwise stated, a masked state removes exactly one legal
398 action; neural adversaries output one removal choice or a no-op. Masks never remove all actions:
399 if removing the selected action would leave an empty set, the original legal set is retained. In two-
400 action games, removing one action naturally forces a singleton; in three- or five-action games, most
401 masks leave multiple alternatives.

Table 6: Q-Learning hyperparameters.

Parameter	Kuhn	Leduc
Learning rate α	0.1	0.1
Exploration ε	0.15	0.15
Discount γ	1.0	1.0
Pre-training episodes	10 000	15 000
Post-mask episodes	10 000	15 000

Table 7: DQN hyperparameters for Leduc-scale experiments.

Parameter	Value
Network (Leduc)	$37 \rightarrow 64 \rightarrow 64 \rightarrow 3$
Network (Leduc-10/20)	$(2N+31) \rightarrow 128 \rightarrow 64 \rightarrow 3$
Learning rate	10^{-3} (Adam)
ε schedule	$1.0 \rightarrow 0.05$ (decay 0.9999)
Replay buffer	20 000–50 000
Batch size	64–128
Target update	every 500–1000 episodes

402 C Hyperparameters

403 D Normalisation Details

404 Normalised performance uses $\text{norm}(r) = (r - r_{\min}) / (r_{\max} - r_{\min})$:

405 Normalisation aids cross-game comparison but obscures absolute exploitability: a normalised .37 in
 406 Leduc ($\Delta r = -3.5$) represents a larger absolute loss than .27 in Kuhn ($\Delta r = -0.9$).

407 E Additional Ablations

408 F Learning Curves Under Attack

409 To confirm that reported results are *asymptotic* (converged) rather than transient, we report reward
 410 in 500-episode windows during post-attack training.

411 **Kuhn + QL (seed 42).** Pre-attack final window: -0.16 . Post-attack trajectory: reward drops to
 412 ~ -0.3 within 2k episodes and oscillates around -0.2 to -0.4 for the remaining 13k episodes. No
 413 recovery trend.

414 **Leduc + QL (seed 42).** Pre-attack: $+0.13$. Post-attack: gradual degradation from 0 to -1.0 over
 415 10k episodes, with oscillations. Final window: -0.40 . The attack takes ~ 5 k episodes to fully
 416 manifest as the adversary’s policy-gradient updates progressively sharpen.

417 **Leduc + NFSP (seed 42).** Pre-attack: -0.30 . Post-attack: initial spike to $+1.0$ (NFSP’s average
 418 strategy briefly exploits the adversary’s early exploration), followed by monotonic decline to -0.51
 419 by episode 15k. The adversary overcomes NFSP’s initial resistance.

420 **Leduc-5 + NFSP (5 seeds).** Pre-attack: -0.06 . Under adversarial masking with 25 outer itera-
 421 tions: progressive degradation from -0.5 (outer 5) to -1.9 (outer 25). Continued training for 15k
 422 additional episodes under the converged mask: reward oscillates between -1.2 and -2.1 with no
 423 recovery trend. The adversary overcomes NFSP’s average-strategy buffer at this scale.

424 These curves confirm that (a) the attack converges and is not a transient phenomenon, (b) victims do
 425 not recover even with $3\times$ extended training under the mask, and (c) the no-recovery property holds
 426 across algorithms and game sizes.

Table 8: Neural NFSP hyperparameters (Leduc-5).

Parameter	Value
Best-response network	$49 \rightarrow 128 \rightarrow 64 \rightarrow 3$
Average-policy network	$49 \rightarrow 128 \rightarrow 64 \rightarrow 3$
BR learning rate	10^{-3}
Average-policy learning rate	10^{-3}
Reservoir buffer	50 000 action samples
Replay buffer	20 000 transitions
Anticipatory parameter η	0.1

Table 9: PPO hyperparameters.

Parameter	Value
Learning rate	0.01
Clip ϵ	0.2
Entropy bonus	0.01
Baseline	Running mean of returns

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [\[Yes\]](#)

Justification: The abstract and introduction state the discrete-action self-play setting, the adversarial action-removal threat model, the empirical scope, and the CACw/CACv mechanism. The Limitations section further bounds the claims to discrete action spaces and empirical scaling trends.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [\[Yes\]](#)

Justification: The paper includes a Limitations section discussing discrete-action scope, local theory, independence assumptions, and remaining full-scale poker/continuous-action gaps.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [\[Yes\]](#)

Justification: Propositions state their assumptions, and the paper gives proof sketches plus an explicit theoretical-scope paragraph. The results are intentionally local/sufficient rather than a global optimal-adversary characterization.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper?

Answer: [\[Yes\]](#)

Justification: The Reproducibility Statement identifies the experiment scripts, fixed seeds, and figure-generation script. Hyperparameters and normalization bounds appear in the appendix.

5. Open access to data and code

Table 10: Adversary hyperparameters (tabular / neural).

Parameter	Tabular	Neural
Architecture	Preference table	$37 \rightarrow 32 \rightarrow 4$ (softmax)
Learning rate	0.01	10^{-3} (Adam)
Outer iterations	20–30	20–25
Inner episodes	500	500
Update rule	REINFORCE	REINFORCE + baseline

Table 11: Normalisation bounds (max single-hand loss/gain).

Game	$r_{\min}$	$r_{\max}$
Kuhn Poker	−2.0	+2.0
Leduc Poker	−13.0	+13.0
Leduc-5 Poker	−17.0	+17.0
Leduc-10 Poker	−21.0	+21.0
Leduc-20 Poker	−41.0	+41.0
Comp. Gridworld	−1.0	+1.0
Resource Collection	−4.0	+4.0

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: The repository includes all environments, agents, adversaries, and scripts. No external datasets are used. For double-blind submission, the repository should be anonymized before linking.

6. Experimental setting/details

Question: Does the paper specify all the training and test details necessary to understand the results?

Answer: [Yes]

Justification: The paper describes self-play, masking, adversary training, action-removal granularity, DQN masking mechanics, and main hyperparameters. Additional details are in Appendix C.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: Main experiments report 95% confidence intervals over seeds where applicable. Some auxiliary diagnostic tables report means and are described as diagnostics.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources needed to reproduce the experiments?

Answer: [Yes]

Justification: Appendix tables report representative sample counts and wall-clock times. Experiments run on CPU with PyTorch used for neural agents and adversaries.

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics?

Answer: [Yes]

Table 12: Mask timing controls in Leduc Q-learning.

Training/evaluation protocol	Victim reward
No mask	$+0.05 \pm 0.03$
Evaluation-only mask	-0.58 ± 0.18
Continued training under mask	-2.65 ± 0.30
Mask-aware training from scratch	-2.71 ± 0.26

Table 13: Threat-model ablation in Leduc Q-learning: private vs. public information adversary.

Adversary information	Victim reward
None	$+0.05 \pm 0.03$
Random mask	-0.98 ± 0.08
Public info only	-1.71 ± 0.58
Private info	-2.30 ± 0.39

488 Justification: The work uses synthetic environments and does not involve human subjects,
 489 private data, or deployed systems. It is framed as robustness/security research with defen-
 490 sive implications.

491 **10. Broader impacts**

492 Question: Does the paper discuss both potential positive societal impacts and negative
 493 societal impacts of the work performed?

494 Answer: [Yes]

495 Justification: The paper discusses capability-removal risks in APIs, actuators, and plat-
 496 forms, and motivates defenses that preserve high-CACv decision capacity. The work could
 497 be misused to identify targeted capability restrictions, but it is presented to motivate robust-
 498 ness testing and defenses.

499 **11. Safeguards**

500 Question: Does the paper describe safeguards that have been put in place for responsible
 501 release of data or models that have a high risk for misuse?

502 Answer: [N/A]

503 Justification: No pretrained models, scraped datasets, or high-risk generative artifacts are
 504 released. The code contains small synthetic RL environments and experiment scripts.

505 **12. Licenses for existing assets**

506 Question: Are the creators or original owners of assets used in the paper properly credited
 507 and are the license and terms of use explicitly mentioned and properly respected?

508 Answer: [Yes]

509 Justification: The code uses standard open-source Python packages and cites related papers.
 510 No external datasets are used. A repository license should be added before broad public
 511 release.

512 **13. New assets**

513 Question: Are new assets introduced in the paper well documented and is the documenta-
 514 tion provided alongside the assets?

515 Answer: [Yes]

516 Justification: New assets are code, synthetic environments, and paper sources. The
 517 README documents layout, dependencies, run commands, scope, and reproducibility.

518 **14. Crowdsourcing and research with human subjects**

519 Question: For crowdsourcing experiments and research with human subjects, does the pa-
 520 per include the full text of instructions and compensation details?

521 Answer: [N/A]

Table 14: CACv-greedy oracle at budget $k = 3$ in Kuhn.

Masking strategy	Victim reward
None	-0.02 ± 0.08
Random	-0.19 ± 0.06
Learned adversary	-0.22 ± 0.02
CACv-greedy oracle	-0.29 ± 0.09

Table 15: Strict matched- L_0 control in Leduc Q-learning. Matched random samples exactly the same number of information states as the learned adversarial mask.

Masking strategy	Effective k	Victim reward
Adversarial	64.8 ± 4.3	-2.32 ± 0.36
Matched random	64.8 ± 4.3	-1.03 ± 0.24

- 522 Justification: No crowdsourcing or human-subject experiments are used.
- 523 15. **Institutional review board (IRB) approvals or equivalent for research with human**
- 524 **subjects**
- 525 Question: Does the paper describe potential participant risks and IRB approvals where
- 526 applicable?
- 527 Answer: [N/A]
- 528 Justification: The work does not involve human subjects.
- 529 16. **Declaration of LLM usage**
- 530 Question: Does the paper describe the usage of LLMs if it is an important, original, or
- 531 non-standard component of the core methods in this research?
- 532 Answer: [N/A]
- 533 Justification: The core methodology, experiments, and analysis do not use LLMs.

Table 16: Precise effective L_0 support diagnostics in Leduc Q-learning.

Mask	Masked states	Seen states	Decision mask rate
Random	69.0	74.0	0.761
Fixed raise	61.4	74.0	0.941
Private adversary	47.6	70.8	0.656
Public adversary	55.0	70.6	0.794

Table 17: Neural and defense ablations.

Ablation	Setting	Victim reward
Separate DQN networks	Leduc + adversary	-1.53 ± 0.48
Defense	Standard training	-2.45 ± 0.60
Defense	Stochastic dropout	-1.82 ± 0.31
Defense	Random mask ensemble	-2.64 ± 0.54

Table 18: Representative compute and sample costs. Wall-clock times are CPU runs on a desktop workstation.

Experiment	Seeds	Episodes per seed	Wall-clock
Leduc DQN scale	5	20k pre + 10k attack	~ 20 min
Leduc-20 DQN scale	5	30k pre + 12.5k attack	~ 17 min
Reviewer ablations	5–10	6k–24k	~ 5 min
Figure generation	n/a	none	< 1 min

Table 19: Budget sweep: victim reward (raw) under adversarial and random masking (Kuhn, 10 seeds, 95% CIs).

k	Adversarial	Random
0	-0.02 ± 0.07	-0.02 ± 0.07
1	-0.21 ± 0.07	-0.02 ± 0.07
2	-0.24 ± 0.07	-0.10 ± 0.06
3	-0.25 ± 0.04	-0.21 ± 0.03
4	-0.28 ± 0.04	-0.41 ± 0.04
5	-0.72 ± 0.21	-0.66 ± 0.01
6	-0.98 ± 0.07	-0.92 ± 0.00

Table 20: Secondary ablations in Kuhn Q-learning.

Ablation	Setting	Victim reward
Self-play	Co-adaptive	-0.975 ± 0.09
Self-play	Fixed opponent	-0.841 ± 0.36
Transfer	Trained on seed 42	-1.032 ± 0.009
Transfer	Retrained per-agent	-0.872 ± 0.134
Parameter sharing	Shared Q-table	-0.948 ± 0.07
Parameter sharing	Separate Q-tables	-0.948 ± 0.07

Table 21: Learned Kuhn mask at full budget.

State	Meaning	Removed	Confidence
0	J, root	BET	0.99
0pb	J, facing bet	PASS	0.99
1	Q, root	PASS	0.95
2	K, root	BET	0.98
2pb	K, facing bet	BET	0.98